Project: Building A Relational Database using Docker, PostgreSQL and DBeaver

Tools Used

- Docker:

Docker (for windows) was downloaded and installed on the desktop for the purpose of containerizing (or housing) the PostgreSQL database to make it easier to manage and deploy the database

- **PostgreSQL:**

PosgreSQL was used because of its robust support for advanced and custom data types, data integrity and ACID compliance amongst other things. These features make it an excellent choice for enterprise level systems.

- **DBeaver:**

Dbeaver is an open-source tool and database management system that facilitates the development and management of the database via an interface. DBeaver was connected to the PostgreSQL container in Docker.

-Command Prompt

The CMD was used to execute the scripts, codes and commands using the **'psql'** command line. The CMD was used for the installation of docker (desktop version), pulling the PostgreSQL image into Docker and containerising it; connecting DBeaver to the Postgresql container in docker and the development of the database.

Project Aim

Aim:

- Build a Retail Database named “retailDB” from scratch using PostgreSQL, Docker and DBeaver
- Employ the use of constraints, keys and referential integrity to ensure data consistency, quality and integrity in the creation of retailDB

Language used for creating retailDB:

- PSQL

Pulled the PostgreSQL image into Docker and created a container for it.

Container name: “some-postgres”

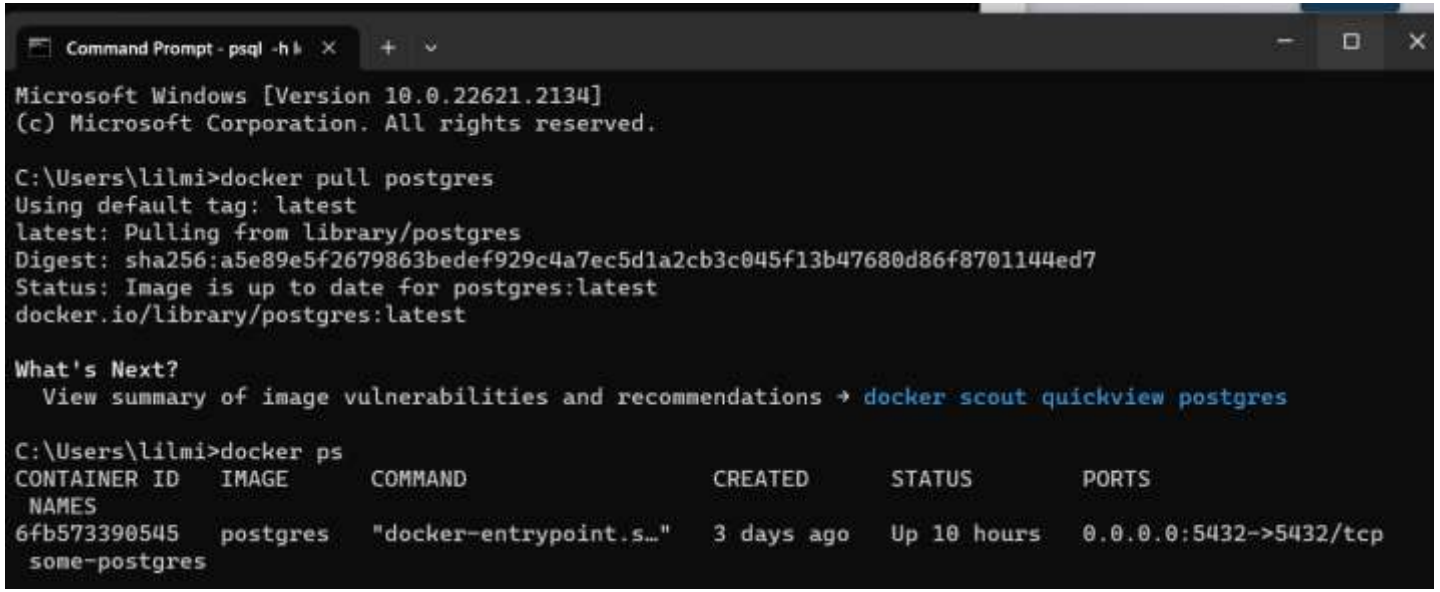

Using the windows command prompt, I created a database in the “some-postgres” container.

Database name= “retailDB”

User = “postgres”

```
C:\Users\lilmi>docker exec some-postgres bash

C:\Users\lilmi>psql -h localhost -U postgres
Password for user postgres:
psql (15.4)
WARNING: Console code page (850) differs from Windows code page (1252)
         8-bit characters might not work correctly. See psql reference
         page "Notes for Windows users" for details.
Type "help" for help.

postgres=# CREATE DATABASE retailDB;
CREATE DATABASE
```

List of databases created by user ‘postgres’ shown below:

```
postgres=# \l
                                                 List of databases
   Name    |  Owner   | Encoding |  Collate   |   Ctype    | ICU Locale | Locale Provider |   Access privileges
-----------+----------+----------+------------+------------+------------+-----------------+-----------------------
 postgres  | postgres | UTF8     | en_US.utf8 | en_US.utf8 |            | libc            |
 retaildb  | postgres | UTF8     | en_US.utf8 | en_US.utf8 |            | libc            |
 template0 | postgres | UTF8     | en_US.utf8 | en_US.utf8 |            | libc            | =c/postgres
+
           |          |          |            |            |            |                 | postgres=CTc/postgres

 template1 | postgres | UTF8     | en_US.utf8 | en_US.utf8 |            | libc            | =c/postgres
+
           |          |          |            |            |            |                 | postgres=CTc/postgres

(4 rows)
```

```
postgres=#
postgres=#
postgres=# \c retaildb;
You are now connected to database "retaildb" as user "postgres".
retaildb=#
```

Snapshot of the newly created “retailDB” database in DBeaver.

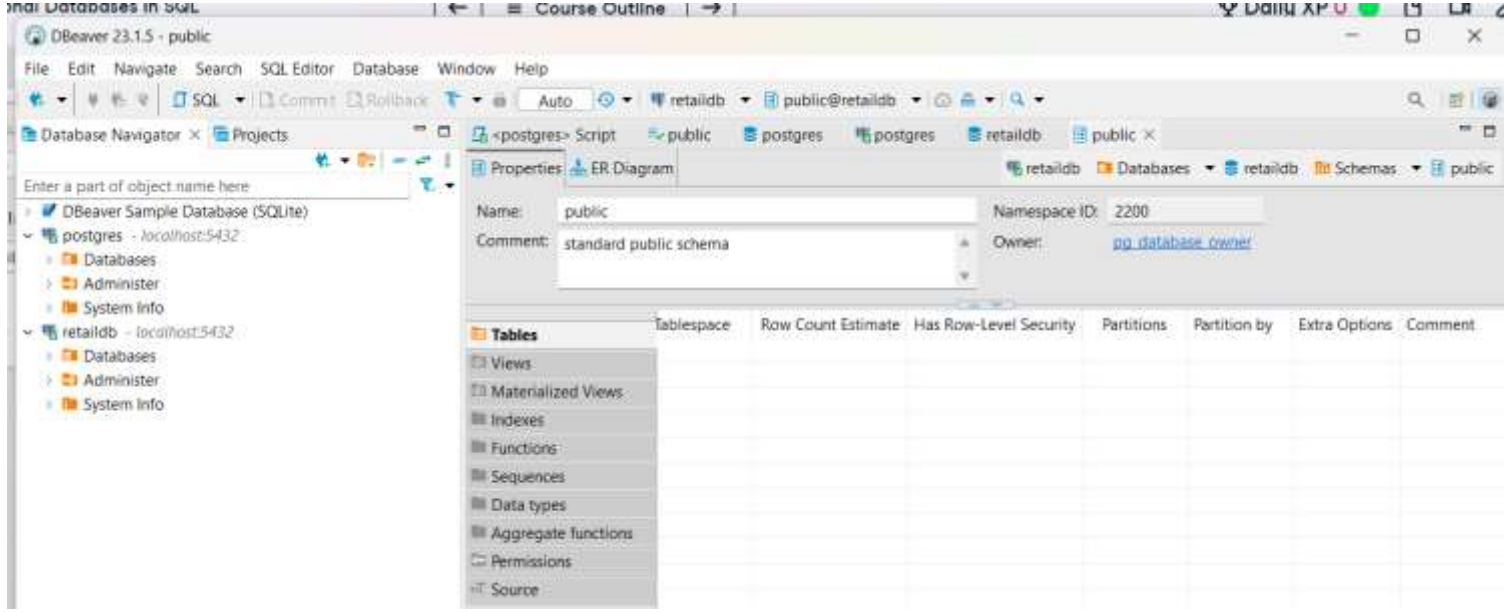

Created the first table in the database called “customer”

Then altered the table to add a primary key constraint using the “Serial” keyword.

- Created the first table ‘customer’ in the retailDB database.

```
postgres=#
postgres=#
postgres=# \c retaildb;
You are now connected to database "retaildb" as user "postgres".
retaildb=# CREATE TABLE customer (firstname TEXT NOT NULL, lastname TEXT NOT NULL
, address VARCHAR(255), cityID INTEGER, phoneNum VARCHAR(255), email VARCHAR(255)
);
CREATE TABLE
```

- Altered the table to add a Primary Key constraint called “customerID”

Altered the table ‘customer’ to add a Primary Key constraint with the auto-increment feature specified by the ‘SERIAL’ keyword, so that each new customer record that is added is automatically given a unique ‘customerID’ that uniquely identifies each row of the customer table.

```
retaildb=#
retaildb=#
retaildb=# ALTER TABLE customer ADD COLUMN customerID serial;
ALTER TABLE
retaildb=#
retaildb=# ALTER TABLE customer ADD CONSTRAINT customer_pkey PRIMARY KEY (customerID);

ALTER TABLE
```

Snapshot of the newly created table in DBeaver

- Customer column with the Primary Key column “customerID”

The columns of the customer table are displayed with their data types. The columns with the ‘NOT NULL’ key constraint are signified by ‘v’. The NOT NULL columns are firstname, lastname and customerid meaning they cannot accept null values.

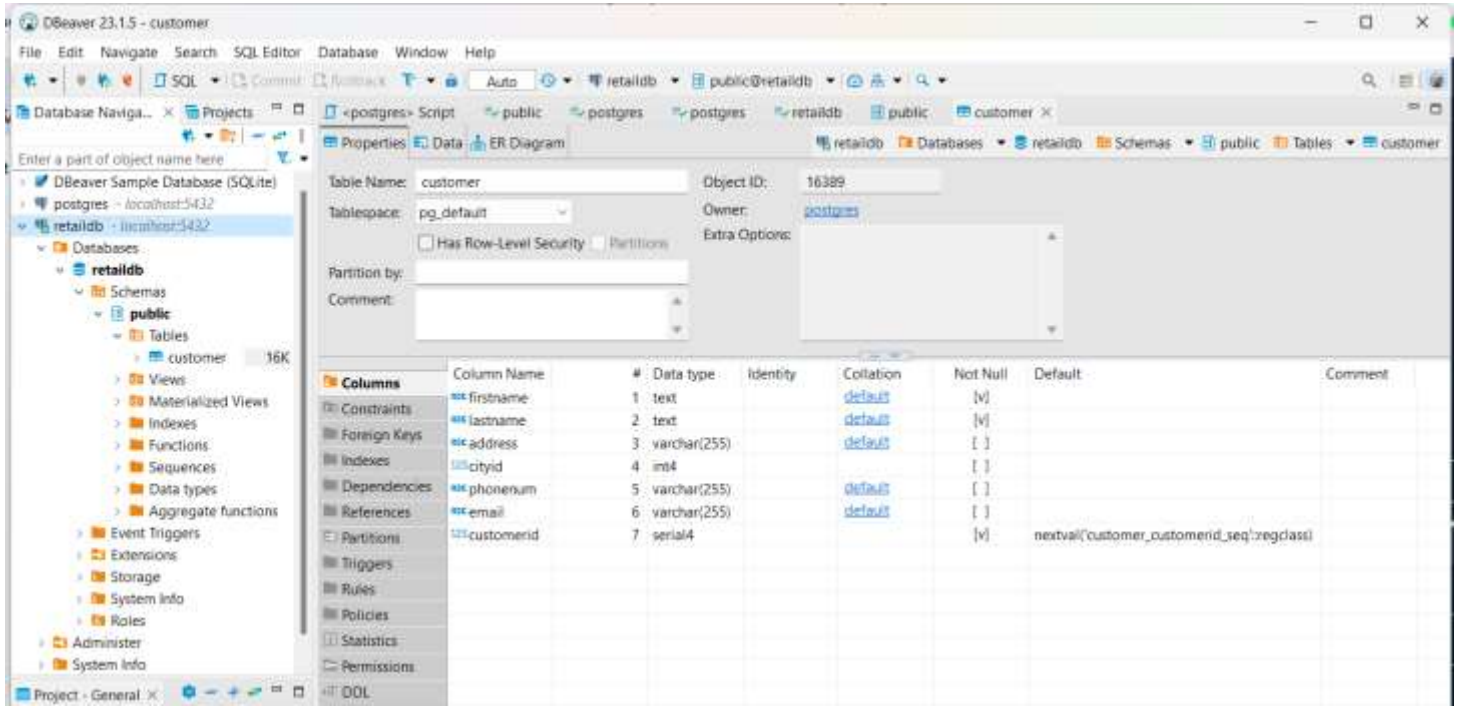

Snapshot showing the creation of 7 more tables (in addition to the customer table) including the fact table “Sales”

- **Creation of 7 new tables using psql in the command prompt**
  - **7 Dimension tables: customer (already created), city, department, employee, orders, product and product category.**
  - **1 Fact Table: Sales**

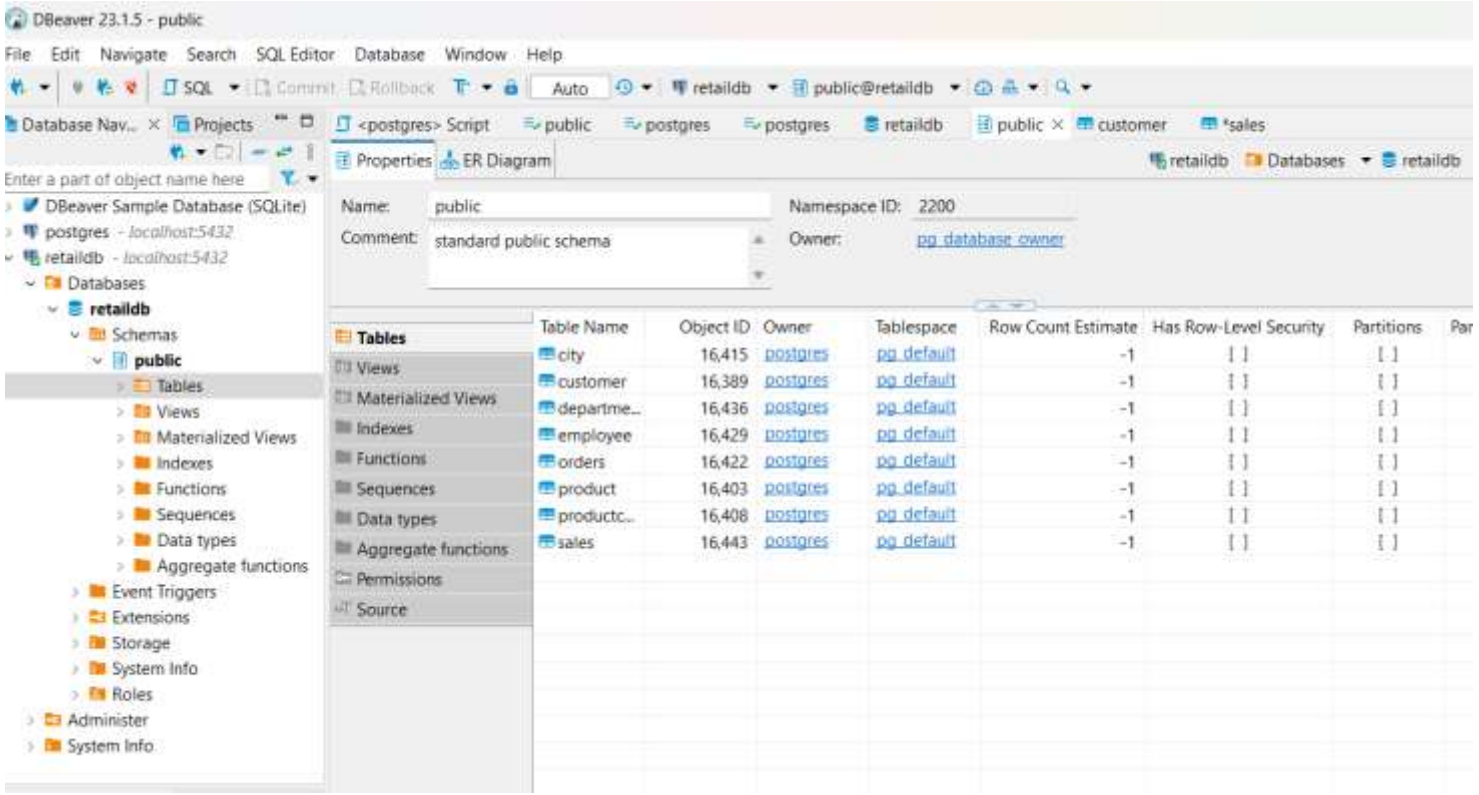

Creation of the Fact table “Sales” in Command Prompt using psql.

- Facts of the Facts table

- The primary key for the Sales table is a combination of the foreign keys being referenced from certain dimension tables (customer, orders, products and employee)

```
retaildb=#
retaildb=#
retaildb=# CREATE TABLE sales (custID INTEGER REFERENCES customer (customerID),
 ordID INTEGER REFERENCES orders (orderID), prodID INTEGER REFERENCES product (
productID), empID INTEGER REFERENCES employee (employeeID), unitsSold INTEGER,
cost NUMERIC (18,2), revenue NUMERIC (25,2));
CREATE TABLE
retaildb=#
```

- Columns of the Sales table in DBeaver

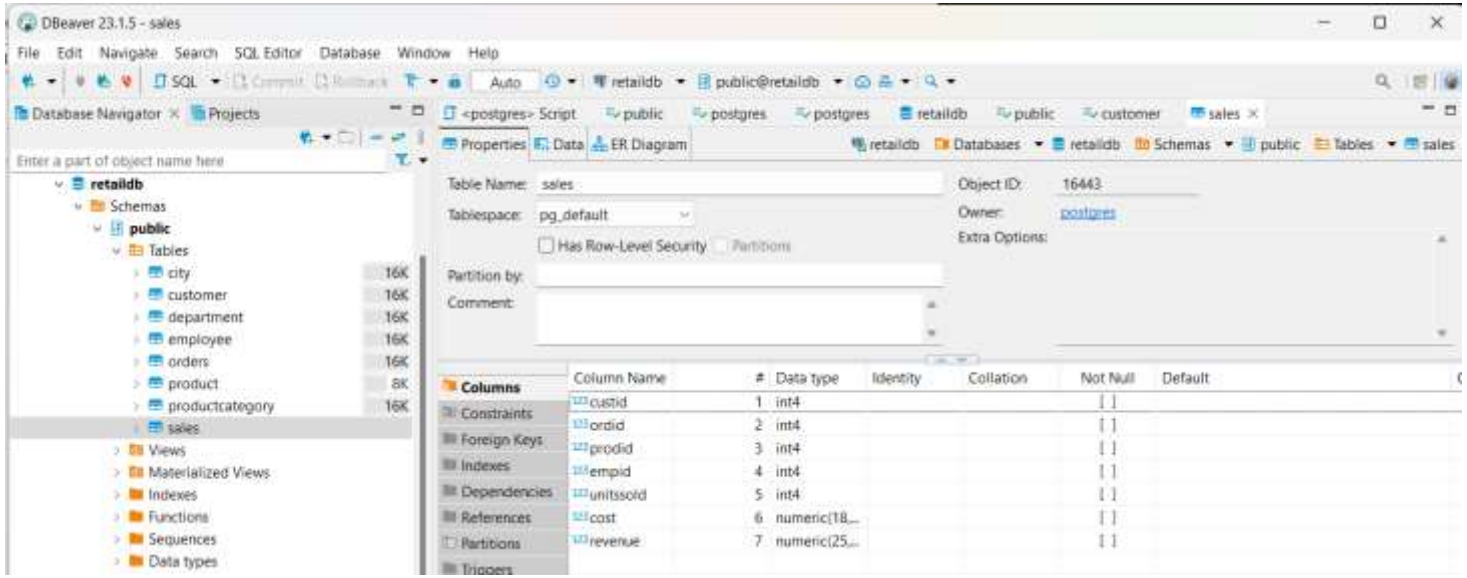

Enforcing Referential Integrity on the columns of the Sales table.

Sales Table showing the Foreign Key columns & Referenced Columns & Associated Tables

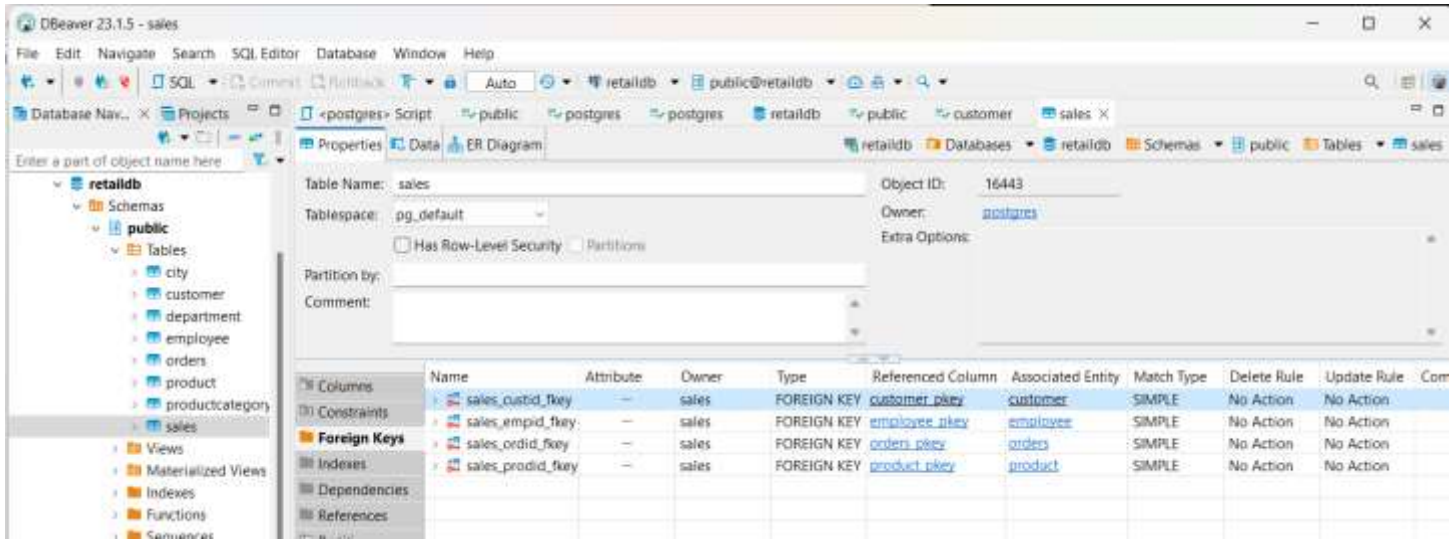

However, foreign key columns by default accept duplicates and null values even if they reference a primary key column in another table. **The problem here does not lie with the duplicate values but rather, the null values,** if null values are allowed in the foreign key columns of the sales table, then it may cause inconsistences in data and ultimately pose a problem for analytics. **Hence, the solution is to set a not null constraint in the foreign key columns of the Sales table using command prompt as shown below;**

```
retaildb=#  ALTER TABLE sales ALTER COLUMN custid SET NOT NULL, ALTER COLUMN ordid
 SET NOT NULL, ALTER COLUMN prodid SET NOT NULL, ALTER COLUMN empid SET NOT NULL;
ALTER TABLE
retaildb=#
```

Enforcing Referential Integrity on the columns of the Sales table by:

1; Adding the 'NOT NULL' constraint to the foreign key columns. This will facilitate a quality relationship between the sales table and the dimension tables thereby promoting data consistency.

NOT NULL constraint on the foreign key columns denoted by 'V'

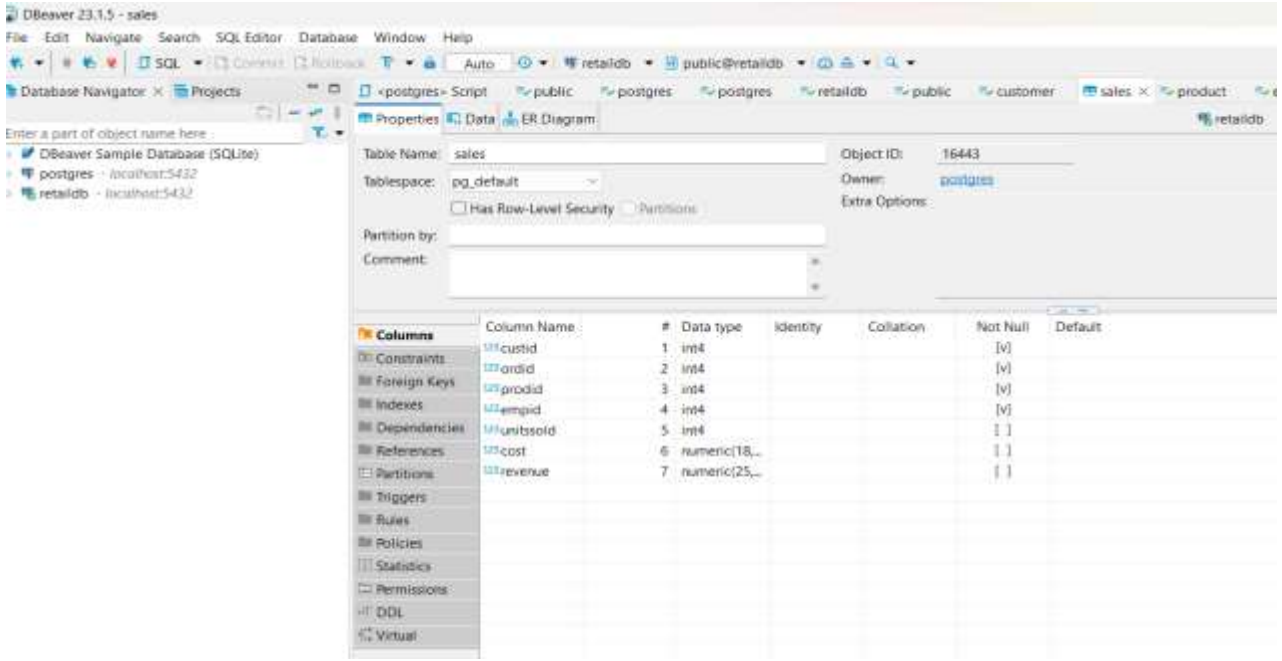

PSQL code showing constraints on the Sales Table

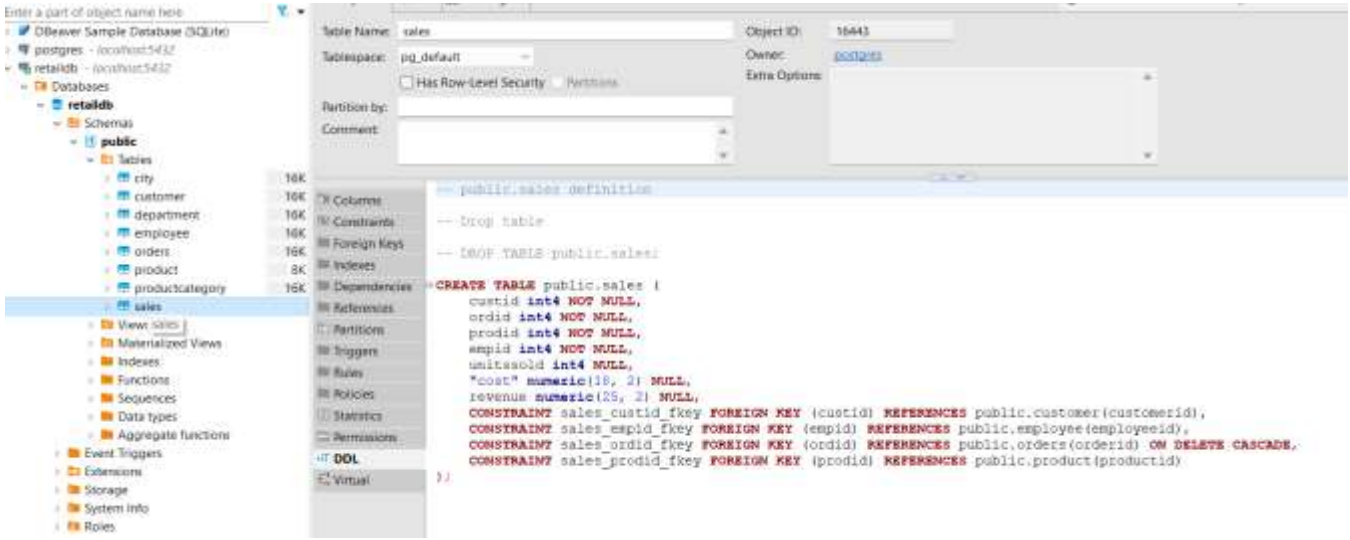

Enforcing Referential Integrity on the columns of the Sales table by

2; Specify the actions to be taken on the Facts table (sales) if a record is deleted from one of the referenced dimensions table.

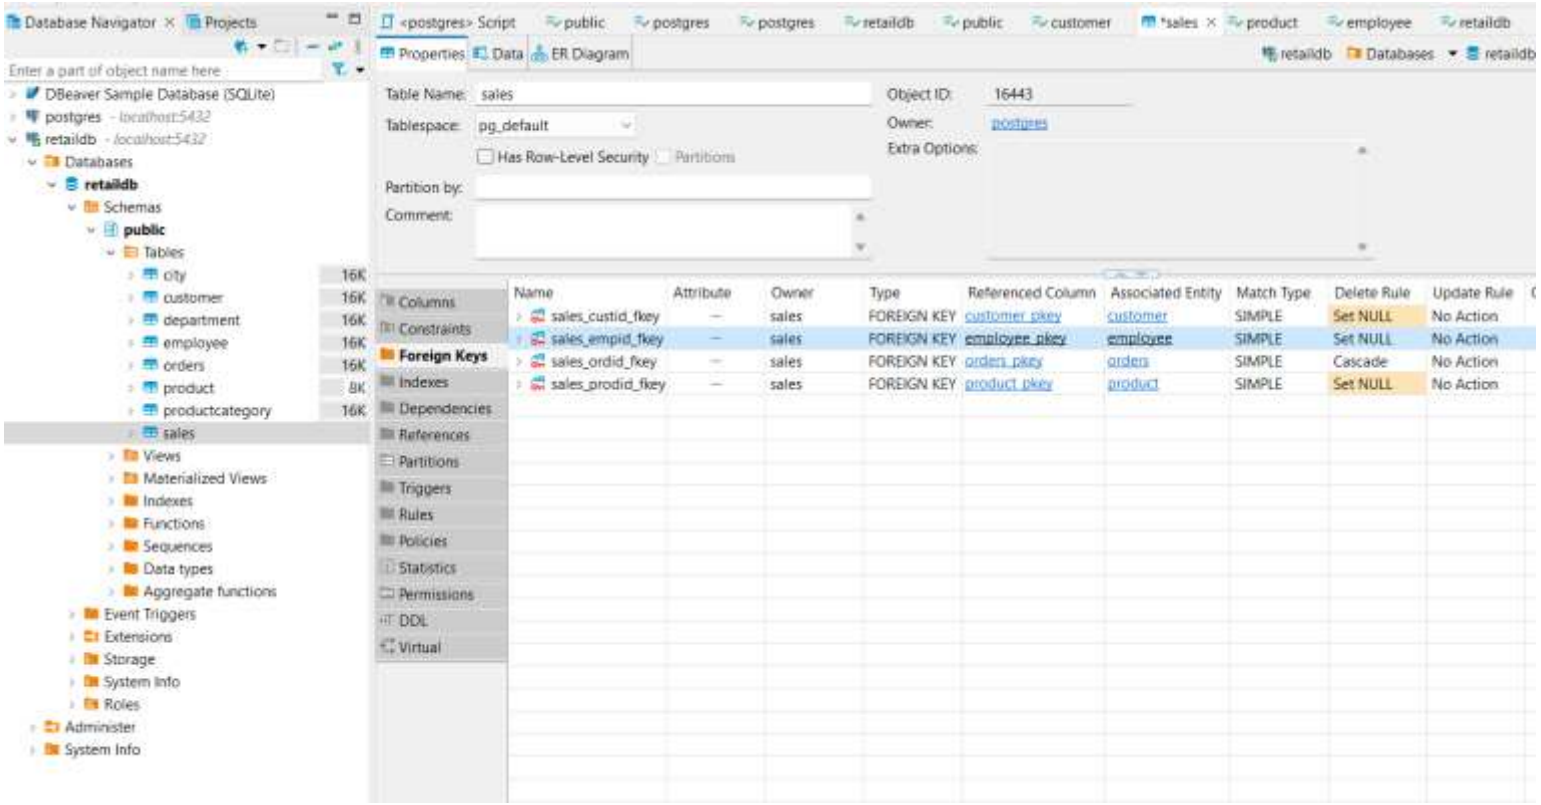

Following the diagram, if a record is deleted from the orders table, then the associated record should be automatically deleted from sales table. For all the other dimension tables, the 'delete rule' is 'Set NULL meaning if a record is deleted in the associated tables, then it should appear as null in the sales table which may lead to some sort of investigation by whoever is handling the data.

Current Entity Relationship Diagram

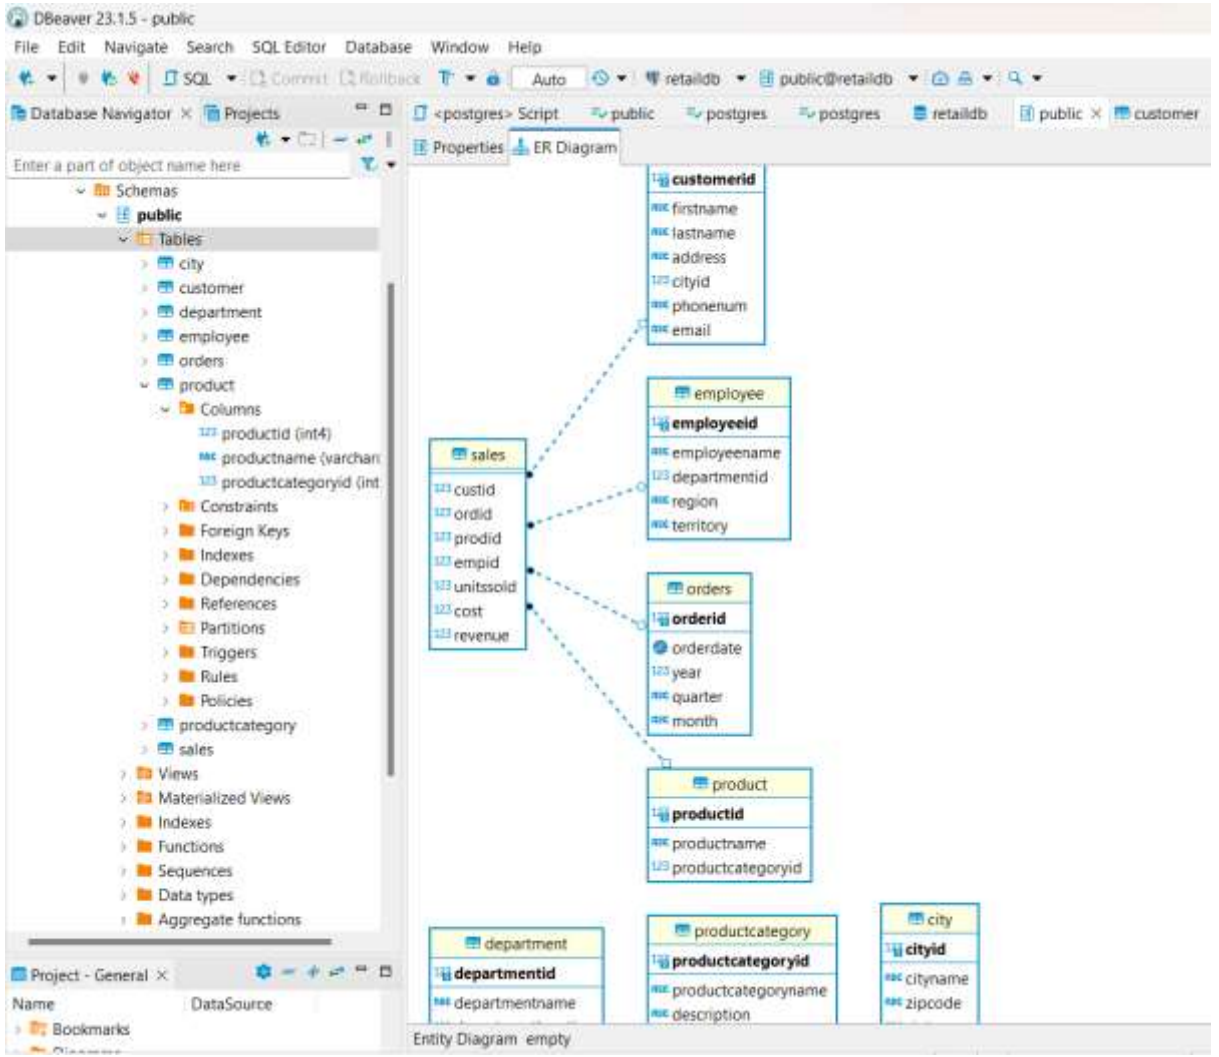

From the diagram, the fact table is connected to only 4 dimension tables which is not the model we're looking for. Hence, a relationship must be established between the facts table and the remaining dimensions table to achieve a 'star schema' or a 'snowflake schema'.

Establishing relationships between the tables by using the ‘REFERENCE’ Keyword

```
retaildb=#
retaildb=#
retaildb=#
retaildb=#
retaildb=#
retaildb=# ALTER TABLE product ADD CONSTRAINT product_fkey FOREIGN KEY (product
CategoryID) REFERENCES productCategory (productCategoryID);
ALTER TABLE
retaildb=# ALTER TABLE employee ADD CONSTRAINT employee_fkey FOREIGN KEY (depar
tmentID) REFERENCES department (departmentID);
ALTER TABLE
retaildb=# ALTER TABLE customer ADD CONSTRAINT customer_fkey FOREIGN KEY (cityI
D) REFERENCES city (cityID);
ALTER TABLE
retaildb=#
```

Remodelled Entity Relationship Diagram

Remodelled Entity Relationship Diagram with a "**Snowflake Schema**"